

# Importing Libraries

Submitted by : Aadarsha Chapagain

StudentID:C0825975

In [2]:

```
import sys
!{sys.executable} -m pip install missingno
```

Collecting missingno

Using cached missingno-0.5.1-py3-none-any.whl (8.7 kB)

Requirement already satisfied: seaborn in s:\anaconda\lib\site-packages (from missingno) (0.11.2)

Requirement already satisfied: matplotlib in s:\anaconda\lib\site-packages (from missingno) (3.4.3)

Requirement already satisfied: numpy in s:\anaconda\lib\site-packages (from missingno) (1.20.3)

Requirement already satisfied: scipy in s:\anaconda\lib\site-packages (from missingno) (1.7.1)

Requirement already satisfied: pyparsing>=2.2.1 in s:\anaconda\lib\site-packages (from matplotlib->missingno) (3.0.4)

Requirement already satisfied: python-dateutil>=2.7 in s:\anaconda\lib\site-packages (from matplotlib->missingno) (2.8.2)

Requirement already satisfied: pillow>=6.2.0 in s:\anaconda\lib\site-packages (from matplotlib->missingno) (8.4.0)

Requirement already satisfied: cycler>=0.10 in s:\anaconda\lib\site-packages (from matplotlib->missingno) (0.10.0)

Requirement already satisfied: kiwisolver>=1.0.1 in s:\anaconda\lib\site-packages (from matplotlib->missingno) (1.3.1)

Requirement already satisfied: six in s:\anaconda\lib\site-packages (from cycler>=0.10->matplotlib->missingno) (1.16.0)

Requirement already satisfied: pandas>=0.23 in s:\anaconda\lib\site-packages (from seaborn->missingno) (1.3.4)

Requirement already satisfied: pytz>=2017.3 in s:\anaconda\lib\site-packages (from pandas>=0.23->seaborn->missingno) (2021.3)

Installing collected packages: missingno

Successfully installed missingno-0.5.1

In [3]:

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import missingno as msno
import seaborn as sns
from scipy import stats
import missingno as msno
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error, accuracy_score
from scipy.stats import shapiro
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.preprocessing import StandardScaler
```

```

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from statsmodels.stats.diagnostic import normal_ad, het_breuschpagan
import statsmodels.api as sm
from sklearn.preprocessing import minmax_scale

%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

```

<https://www.kaggle.com/uciml/red-wine-quality-cortez-et-al-2009>

## Datasets

In [4]: `df = pd.read_csv('winequality-red.csv')`

In [5]: `df.head()`

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5

In [6]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
0   fixed acidity       1599 non-null   float64

```

```

1  volatile acidity      1599 non-null  float64
2  citric acid           1599 non-null  float64
3  residual sugar        1599 non-null  float64
4  chlorides             1599 non-null  float64
5  free sulfur dioxide   1599 non-null  float64
6  total sulfur dioxide  1599 non-null  float64
7  density               1599 non-null  float64
8  pH                   1599 non-null  float64
9  sulphates             1599 non-null  float64
10 alcohol               1599 non-null  float64
11 quality               1599 non-null  int64

```

```
dtypes: float64(11), int64(1)
```

```
memory usage: 150.0 KB
```

fixed acidity: most acids involved with wine or fixed or nonvolatile (do not evaporate readily)

volatile acidity: the amount of acetic acid in wine, which at too high of levels can lead to an unpleasant, vinegar taste

citric acid: found in small quantities, citric acid can add 'freshness' and flavor to wines

residual sugar: the amount of sugar remaining after fermentation stops, it's rare to find wines with less than 1 gram/liter and wines with greater than 45 grams/liter are considered sweet

chlorides: the amount of salt in the wine

free sulfur dioxide: the free form of S02 exists in equilibrium between molecular S02 (as a dissolved gas) and bisulfite ion; it prevents microbial growth and the oxidation of wine

total sulfur dioxide: amount of free and bound forms of S02; in low concentrations, S02 is mostly undetectable in wine, but at free S02 concentrations over 50 ppm, S02 becomes evident in the nose and taste of wine

density: the density of water is close to that of water depending on the percent alcohol and sugar content

pH: describes how acidic or basic a wine is on a scale from 0 (very acidic) to 14 (very basic); most wines are between 3-4 on the pH scale

sulphates: a wine additive which can contribute to sulfur dioxide gas (S02) levels, wich acts as an antimicrobial and antioxidant

quality: score between 0 and 10


alcohol: the percent alcohol content of the wine - output varaible

## Missing Data

```
In [7]: msno.matrix(df)
```

```
<AxesSubplot:>
```

Out[7]:



## Data visulization

In [8]:

```
def plotBoxplot(data):
    fig, axes = plt.subplots(ncols=3, nrows=4, figsize=(15,15))
    fig.tight_layout(pad=4.0)

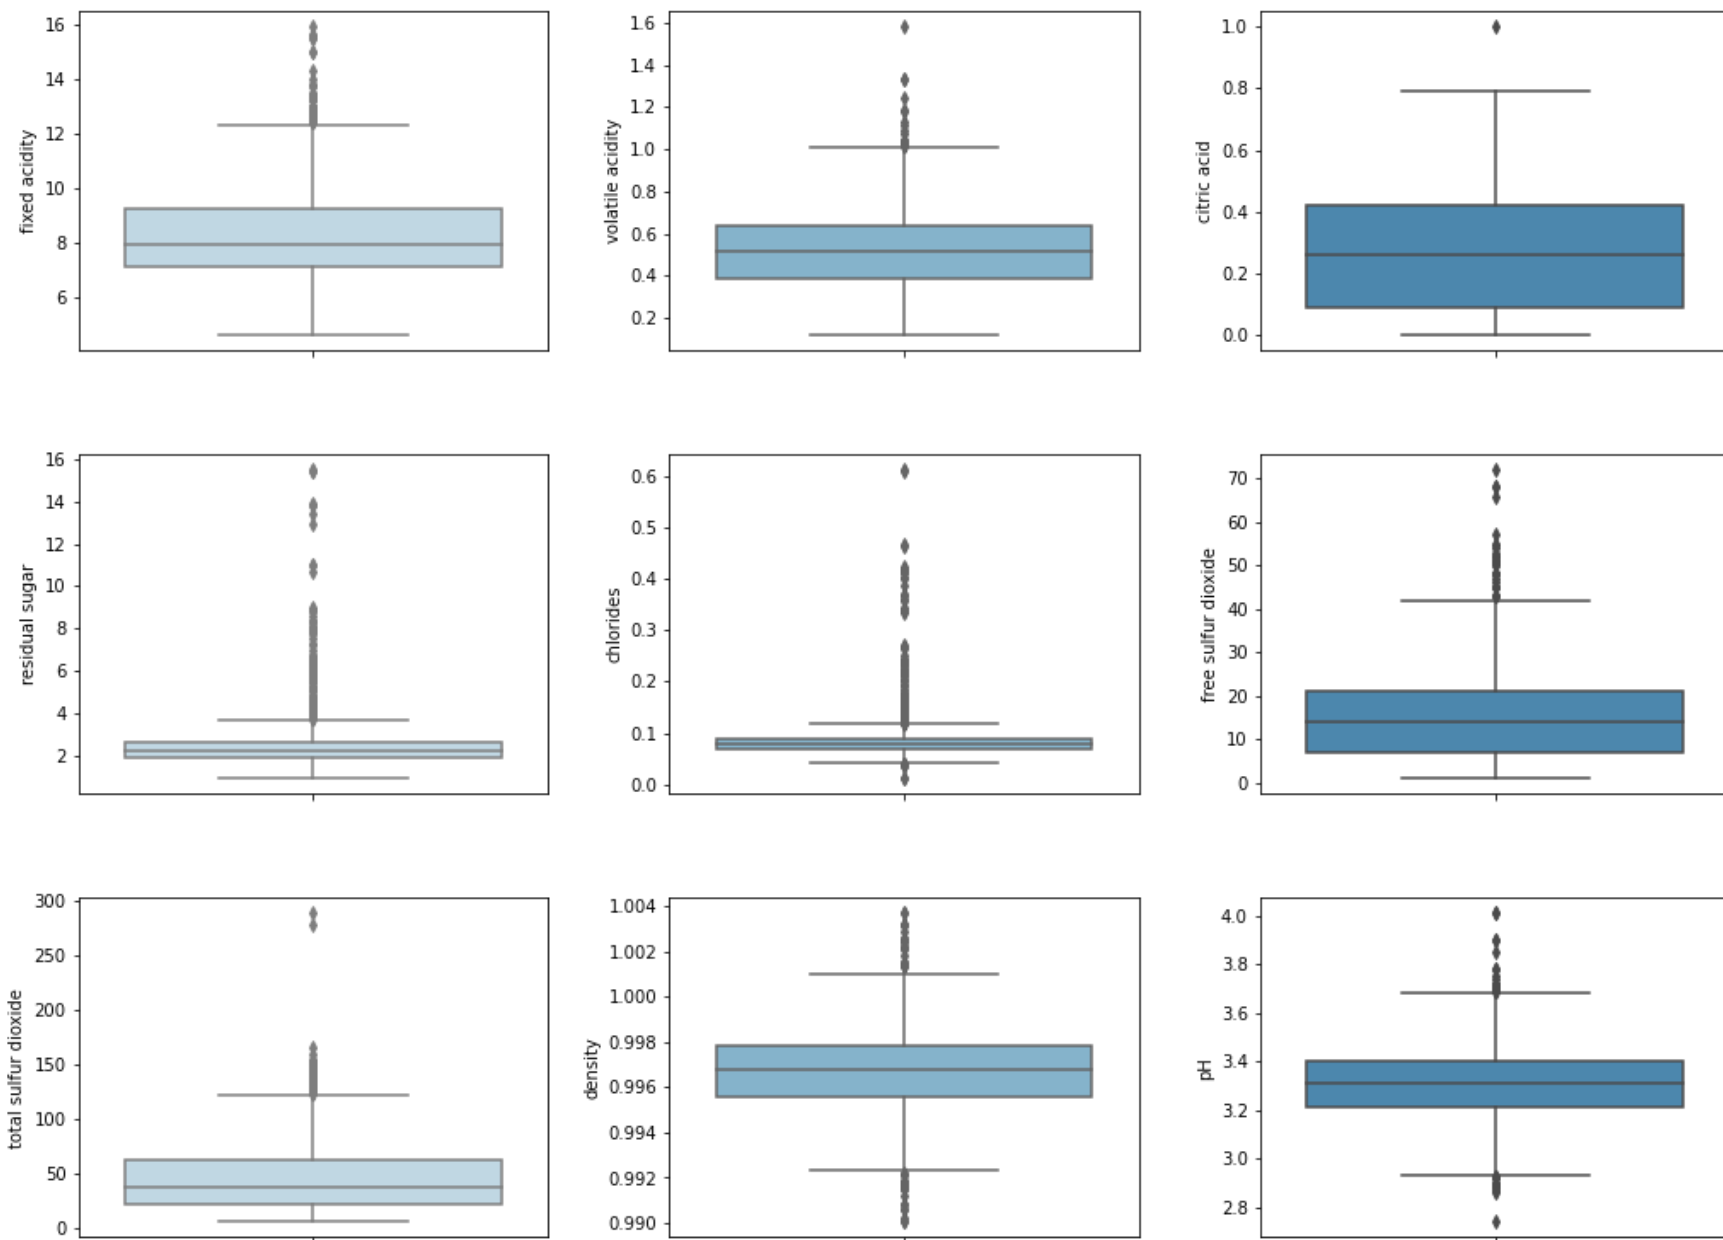
    col = 0
    row = 0
    colors = ['#bad9e9', '#7ab6d6', '#3c8abd']


    for i, column in enumerate(data.columns):
        sns.boxplot(y=column, data=data, ax=axes[row][col], color=colors[col])

        if (i + 1) % 3 == 0:
            row += 1
```

```
col = 0  
else:  
col += 1
```

```
plotBoxplot(df)
```






No Multicollinearity among Independent Variables


```
In [10]: plt.figure(figsize=(25, 10))
sns.heatmap(df.loc[:, df.columns != 'alcohol'].corr(), annot=True, fmt='.2f')
```

Out[10]: <AxesSubplot:>




```
In [11]: #Here we see that fixed acidity does not give any specification to classify the quality.  
fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'fixed acidity', data = df)
```

```
Out[11]: <AxesSubplot:xlabel='quality', ylabel='fixed acidity'>
```



```
In [12]: #Here we see that its quite a downing trend in the volatile acidity as we go higher the quality  
fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'volatile acidity', data = df)
```


```
Out[12]: <AxesSubplot:xlabel='quality', ylabel='volatile acidity'>
```



```
In [13]: #Composition of citric acid go higher as we go higher in the quality of the wine  
fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'citric acid', data = df)
```


```
Out[13]: <AxesSubplot:xlabel='quality', ylabel='citric acid'>
```






```
In [14]: fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'residual sugar', data = df)
```

```
Out[14]: <AxesSubplot:xlabel='quality', ylabel='residual sugar'>
```




```
In [15]: #Composition of chloride also go down as we go higher in the quality of the wine  
fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'chlorides', data = df)
```

```
Out[15]: <AxesSubplot:xlabel='quality', ylabel='chlorides'>
```




```
In [16]: fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'free sulfur dioxide', data = df)
```

```
Out[16]: <AxesSubplot:xlabel='quality', ylabel='free sulfur dioxide'>
```




```
In [17]: fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'total sulfur dioxide', data = df)
```

```
Out[17]: <AxesSubplot:xlabel='quality', ylabel='total sulfur dioxide'>
```




```
In [18]: #Sulphates level goes higher with the quality of wine  
fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'sulphates', data = df)
```

```
Out[18]: <AxesSubplot:xlabel='quality', ylabel='sulphates'>
```




```
In [19]: #Alcohol Level also goes higher as te quality of wine increases  
fig = plt.figure(figsize = (10,6))  
sns.barplot(x = 'quality', y = 'alcohol', data = df)
```

```
Out[19]: <AxesSubplot:xlabel='quality', ylabel='alcohol'>
```



```
In [20]: sns.countplot(df['quality'])
```

```
Out[20]: <AxesSubplot:xlabel='quality', ylabel='count'>
```



# Features and Labels

```
In [21]: #create tmp train/test split for assumptions test
X = df.drop(['alcohol'], axis=1)
y = df['alcohol']
```

```
In [19]: print(X)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.700	0.00	1.9	0.076	
1	7.8	0.880	0.00	2.6	0.098	
2	7.8	0.760	0.04	2.3	0.092	
3	11.2	0.280	0.56	1.9	0.075	
4	7.4	0.700	0.00	1.9	0.076	
...	...	...	...	...	...	
1594	6.2	0.600	0.08	2.0	0.090	
1595	5.9	0.550	0.10	2.2	0.062	
1596	6.3	0.510	0.13	2.3	0.076	
1597	5.9	0.645	0.12	2.0	0.075	
1598	6.0	0.310	0.47	3.6	0.067	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.99780	3.51	0.56	
1	25.0	67.0	0.99680	3.20	0.68	
2	15.0	54.0	0.99700	3.26	0.65	
3	17.0	60.0	0.99800	3.16	0.58	
4	11.0	34.0	0.99780	3.51	0.56	
...	...	...	...	...	...	
1594	32.0	44.0	0.99490	3.45	0.58	
1595	39.0	51.0	0.99512	3.52	0.76	
1596	29.0	40.0	0.99574	3.42	0.75	
1597	32.0	44.0	0.99547	3.57	0.71	
1598	18.0	42.0	0.99549	3.39	0.66	

	quality
0	5
1	5
2	5
3	6
4	5
...	...
1594	5



```
1595      6
1596      6
1597      5
1598      6
```

```
[1599 rows x 11 columns]
```

In [20]:

```
print(y)
```

```
0      9.4
1      9.8
2      9.8
3      9.8
4      9.4
```

```
...
1594   10.5
1595   11.2
1596   11.0
1597   10.2
1598   11.0
```

```
Name: alcohol, Length: 1599, dtype: float64
```

## Splitting the dataset

In [22]:

```
x_train, x_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=50)
```

In [22]:

```
print(x_train)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
940	9.6	0.330	0.52	2.2	0.074	
1287	8.0	0.600	0.08	2.6	0.056	
1397	7.3	0.590	0.26	2.0	0.080	
356	11.5	0.410	0.52	3.0	0.080	
226	8.9	0.590	0.50	2.0	0.337	
...	...	...	...	...	...	
70	7.7	0.630	0.08	1.9	0.076	
132	5.6	0.500	0.09	2.3	0.049	
1313	7.0	0.360	0.21	2.3	0.086	
109	8.1	0.785	0.52	2.0	0.122	
1504	7.5	0.380	0.57	2.3	0.106	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
940	13.0	25.0	0.99509	3.36	0.76
1287	3.0	7.0	0.99286	3.22	0.37
1397	17.0	104.0	0.99584	3.28	0.52
356	29.0	55.0	1.00010	3.26	0.88
226	27.0	81.0	0.99640	3.04	1.61
...	...	...	...	...	...
70	15.0	27.0	0.99670	3.32	0.54
132	17.0	99.0	0.99370	3.63	0.63
1313	20.0	65.0	0.99558	3.40	0.54
109	37.0	153.0	0.99690	3.21	0.69
1504	5.0	12.0	0.99605	3.36	0.55

	quality
940	7
1287	5
1397	5
356	5
226	6
...	...
70	6
132	5
1313	6
109	5
1504	6

[1119 rows x 11 columns]

In [23]:

```
print(y_train)
```

940	12.4
1287	13.0
1397	9.9
356	11.0
226	9.5
...	...
70	9.5
132	13.0
1313	10.1
109	9.3
1504	11.4

Name: alcohol, Length: 1119, dtype: float64

In [24]:

```
print(x_test)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
453	10.4	0.33	0.63	2.80	0.084	
1415	6.2	0.58	0.00	1.60	0.065	
1242	9.0	0.40	0.41	2.00	0.058	
885	8.9	0.75	0.14	2.50	0.086	
488	11.6	0.32	0.55	2.80	0.081	
...	...	...	...	...	...	
34	5.2	0.32	0.25	1.80	0.103	
1493	7.7	0.54	0.26	1.90	0.089	
501	10.4	0.44	0.73	6.55	0.074	
1464	6.8	0.59	0.10	1.70	0.063	
911	9.1	0.28	0.46	9.00	0.114	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
453	5.0	22.0	0.99980	3.26	0.74	
1415	8.0	18.0	0.99660	3.56	0.84	
1242	15.0	40.0	0.99414	3.22	0.60	
885	9.0	30.0	0.99824	3.34	0.64	
488	35.0	67.0	1.00020	3.32	0.92	
...	...	...	...	...	...	
34	13.0	50.0	0.99570	3.38	0.55	
1493	23.0	147.0	0.99636	3.26	0.59	
501	38.0	76.0	0.99900	3.17	0.85	
1464	34.0	53.0	0.99580	3.41	0.67	
911	3.0	9.0	0.99901	3.18	0.60	

	quality
453	7
1415	5
1242	6
885	5
488	7
...	...
34	5
1493	5
501	7
1464	5
911	6

[480 rows x 11 columns]

```
In [24]: print(y_test)
```

453 11.2

```
1415      9.4
1242     12.2
885      10.5
488      10.8
...
34       9.2
1493     9.7
501     12.0
1464     9.7
911     10.9
```

Name: alcohol, Length: 480, dtype: float64

```
In [25]: print(x_train.shape)
        print(y_train.shape)
```

```
(1119, 11)
(1119,)
```

```
In [26]: sc = StandardScaler()
        X_train = sc.fit_transform(x_train)
        X_test = sc.transform(x_test)
```

## Model Setup

```
In [27]: from sklearn.linear_model import LinearRegression
```

```
In [28]: model = LinearRegression()
```

## Training

```
In [30]: model.fit(x_train, y_train)
```

```
Out[30]: LinearRegression()
```

```
In [31]: model.coef_
```

```
Out[31]: array([ 5.05933116e-01,  7.02590521e-01,  8.23700660e-01,  2.72656929e-01,  
        -5.33212379e-01, -5.68110422e-03, -5.96585162e-04, -5.87493452e+02,  
         3.78523263e+00,  8.59410226e-01,  2.47684402e-01])
```

```
In [32]: model.intercept_
```

```
Out[32]: 576.1956939510893
```

## Evaluating Model

```
In [33]: print(model.score(X, y))
```

```
0.6895215542808213
```

## Predicting on Test data


```
In [34]: y_pred = model.predict(x_test)
```

```
In [35]: r2_score(y_test, y_pred)
```

```
Out[35]: 0.6648714487925931
```

## Visualisation

```
In [36]: #plot the actual vs predicted values  
sns.regplot(y_test, y_pred, line_kws={'color':'red'}, ci=None)  
  
plt.xlabel('Actual')  
plt.ylabel('Predictions')  
plt.title('Prediction vs Actual')  
  
plt.show()
```



## References

<https://learningwithdata.com/posts/tylerfolkman/the-ultimate-guide-to-linear-regression/>  
<https://www.analyticsvidhya.com/blog/2021/05/all-you-need-to-know-about-your-first-machine-learning-model-linear-regression/>  
<https://towardsdatascience.com/the-complete-guide-to-linear-regression-in-python-3d3f8f06bf8>  
<https://www.keboola.com/blog/linear-regression-machine-learning>  
[https://rstudio-pubs-static.s3.amazonaws.com/57835\\_c4ace81da9dc45438ad0c286bcbb4224.html](https://rstudio-pubs-static.s3.amazonaws.com/57835_c4ace81da9dc45438ad0c286bcbb4224.html)  
<https://www.itl.nist.gov/div898/handbook/prc/section1/prc16.htm>  
<https://www.kaggle.com/nareshbhat/outlier-the-silent-killer>  
<https://towardsdatascience.com/ways-to-detect-and-remove-the-outliers-404d16608dba>  
<https://community.gooddata.com/metrics-and-maql-kb-articles-43/normality-testing-skewness-and-kurtosis-241>  
<https://towardsdatascience.com/methods-for-normality-test-with-application-in-python-bb91b49ed0f5>  
<https://medium.com/@TheDataGyan/day-8-data-transformation-skewness-normalization-and-much-more-4c144d370e55>  
<https://machinelearningmastery.com/standardscaler-and-minmaxscaler-transforms-in-python/>  
<https://medium.com/@TheDataGyan/day-8-data-transformation-skewness-normalization-and-much-more-4c144d370e55>  
<https://jyotiyadav99111.medium.com/statistics-how-should-i-interpret-results-of-ols-3bde1ebeec01>