

# chapter5

March 2, 2022

## 1 Chapter 5 - Basic Math and Statistics

### 1.1 Segment 1 - Using NumPy to perform arithmetic operations on data

```
[1]: import numpy as np
      from numpy.random import randn
```

```
[2]: np.set_printoptions(precision=2)
```

### 1.2 Creating arrays

#### 1.2.1 Creating arrays using a list

```
[3]: a = np.array([1,2,3,4,5,6])
      a
```

```
[3]: array([1, 2, 3, 4, 5, 6])
```

```
[4]: b = np.array([[10,20,20],[40,50,60]])
      b
```

```
[4]: array([[10, 20, 20],
           [40, 50, 60]])
```

#### 1.2.2 Creating arrays via assignment

```
[5]: np.random.seed(25)
      c = 36*np.random.randn(6)
      c
```

```
[5]: array([ 8.22, 36.97, -30.23, -21.28, -34.45, -8.  ])
```

```
[6]: d = np.arange(1, 35)
      d
```

```
[6]: array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17,
          18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34])
```

### 1.3 Performing arithmetic on arrays

```
[7]: a*10
```

```
[7]: array([10, 20, 30, 40, 50, 60])
```

```
[8]: c + a
```

```
[8]: array([ 9.22, 38.97, -27.23, -17.28, -29.45, -2.  ])
```

```
[9]: c-a
```

```
[9]: array([ 7.22, 34.97, -33.23, -25.28, -39.45, -14.  ])
```

```
[10]: c*a
```

```
[10]: array([ 8.22, 73.94, -90.68, -85.13, -172.24, -48.02])
```

```
[11]: c/a
```

```
[11]: array([ 8.22, 18.48, -10.08, -5.32, -6.89, -1.33])
```

#### 1.3.1 Multiplying matrices and basic linear algebra

```
[12]: aa = np.array([[2.,4.,6.],[1.,3.,5.],[10.,20.,30.]])  
aa
```

```
[12]: array([[ 2.,  4.,  6.],  
            [ 1.,  3.,  5.],  
            [10., 20., 30.]])
```

```
[13]: bb = np.array([[0.,1.,2.],[3.,4.,5.],[6.,7.,8.]])  
bb
```

```
[13]: array([[0., 1., 2.],  
            [3., 4., 5.],  
            [6., 7., 8.]])
```

```
[14]: aa*bb
```

```
[14]: array([[ 0.,  4., 12.],  
            [ 3., 12., 25.],  
            [60., 140., 240.]])
```

```
[15]: np.dot(aa,bb)
```

```
[15]: array([[ 48.,  60.,  72.],  
            [ 39.,  48.,  57.],  
            [240., 300., 360.]])
```

## 1.4 Segment 2 - Multiplying matrices and basic linear algebra

### 1.5 Multiplying matrices and basic linear algebra

```
[16]: aa = np.array([[2.,4.,6.],[1.,3.,5.],[10.,20.,30.]])
      aa
```

```
[16]: array([[ 2.,  4.,  6.],
           [ 1.,  3.,  5.],
           [10., 20., 30.]])
```

```
[17]: bb = np.array([[0.,1.,2.],[3.,4.,5.],[6.,7.,8.]])
      bb
```

```
[17]: array([[0., 1., 2.],
           [3., 4., 5.],
           [6., 7., 8.]])
```

```
[18]: aa*bb
```

```
[18]: array([[ 0.,  4., 12.],
           [ 3., 12., 25.],
           [60., 140., 240.]])
```

```
[19]: np.dot(aa,bb)
```

```
[19]: array([[ 48.,  60.,  72.],
           [ 39.,  48.,  57.],
           [240., 300., 360.]])
```

## 1.6 Segment 3 - Generating summary statistics using pandas and scipy

```
[22]: import numpy as np
      import pandas as pd
      from pandas import Series, DataFrame

      import scipy
      from scipy import stats
```

```
[23]: address = './Data/mtcars.csv'

      cars = pd.read_csv(address)
      cars.columns =
      → ['car_names', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am', 'gear', 'carb']

      cars.head()
```

```
[23]:
```

	car_names	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	\
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	

	carb
0	4
1	4
2	1
3	1
4	2

```
[24]: address = './Data/mtcars.csv'

cars = pd.read_csv(address)
cars.columns = \
    ['car_names', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am', 'gear', 'carb']

cars.head()
```

```
[24]:
```

	car_names	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	\
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	

	carb
0	4
1	4
2	1
3	1
4	2

### 1.6.1 Looking at summary statistics that describe a variable's numeric values

```
[26]: cars.sum()
```

```
[26]: car_names      Mazda RX4Mazda RX4 WagDatsun 710Hornet 4 Drive...
mpg                                642.9
cyl                                198
disp                             7383.1
hp                                4694
drat                             115.09
wt                                102.952
```

```

qsec          571.16
vs            14
am            13
gear          118
carb          90
dtype: object

```

```
[27]: cars.sum(axis=1)
```

```

[27]: 0      328.980
      1      329.795
      2      259.580
      3      426.135
      4      590.310
      5      385.540
      6      656.920
      7      270.980
      8      299.570
      9      350.460
     10      349.660
     11      510.740
     12      511.500
     13      509.850
     14      728.560
     15      726.644
     16      725.695
     17      213.850
     18      195.165
     19      206.955
     20      273.775
     21      519.650
     22      506.085
     23      646.280
     24      631.175
     25      208.215
     26      272.570
     27      273.683
     28      670.690
     29      379.590
     30      694.710
     31      288.890
dtype: float64

```

```
[28]: cars.median()
```

```

[28]: mpg      19.200
      cyl       6.000

```

```
disp    196.300
hp      123.000
drat     3.695
wt       3.325
qsec    17.710
vs       0.000
am       0.000
gear     4.000
carb     2.000
dtype: float64
```

```
[29]: cars.mean()
```

```
[29]: mpg      20.090625
      cyl      6.187500
      disp   230.721875
      hp     146.687500
      drat    3.596563
      wt      3.217250
      qsec   17.848750
      vs      0.437500
      am      0.406250
      gear    3.687500
      carb    2.812500
      dtype: float64
```

```
[30]: cars.max()
```

```
[30]: car_names  Volvo 142E
      mpg      33.9
      cyl       8
      disp   472.0
      hp      335
      drat    4.93
      wt     5.424
      qsec   22.9
      vs       1
      am       1
      gear     5
      carb     8
      dtype: object
```

```
[31]: mpg = cars.mpg
      mpg.idxmax()
```

```
[31]: 19
```

## 1.6.2 Looking at summary statistics that describe variable distribution

```
[33]: cars.std()
```

```
[33]: mpg      6.026948
      cyl      1.785922
      disp    123.938694
      hp      68.562868
      drat     0.534679
      wt      0.978457
      qsec     1.786943
      vs      0.504016
      am      0.498991
      gear     0.737804
      carb     1.615200
      dtype: float64
```

```
[34]: cars.var()
```

```
[34]: mpg      36.324103
      cyl      3.189516
      disp    15360.799829
      hp      4700.866935
      drat     0.285881
      wt      0.957379
      qsec     3.193166
      vs      0.254032
      am      0.248992
      gear     0.544355
      carb     2.608871
      dtype: float64
```

```
[35]: gear = cars.gear
      gear.value_counts()
```

```
[35]: 3    15
      4    12
      5     5
      Name: gear, dtype: int64
```

```
[36]: cars.describe()
```

```
[36]:
```

	mpg	cyl	disp	hp	drat	wt \
count	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000
mean	20.090625	6.187500	230.721875	146.687500	3.596563	3.217250
std	6.026948	1.785922	123.938694	68.562868	0.534679	0.978457
min	10.400000	4.000000	71.100000	52.000000	2.760000	1.513000
25%	15.425000	4.000000	120.825000	96.500000	3.080000	2.581250

50%	19.200000	6.000000	196.300000	123.000000	3.695000	3.325000
75%	22.800000	8.000000	326.000000	180.000000	3.920000	3.610000
max	33.900000	8.000000	472.000000	335.000000	4.930000	5.424000

	qsec	vs	am	gear	carb
count	32.000000	32.000000	32.000000	32.000000	32.0000
mean	17.848750	0.437500	0.406250	3.687500	2.8125
std	1.786943	0.504016	0.498991	0.737804	1.6152
min	14.500000	0.000000	0.000000	3.000000	1.0000
25%	16.892500	0.000000	0.000000	3.000000	2.0000
50%	17.710000	0.000000	0.000000	4.000000	2.0000
75%	18.900000	1.000000	1.000000	4.000000	4.0000
max	22.900000	1.000000	1.000000	5.000000	8.0000

## 1.7 Segment 4 - Summarizing categorical data using pandas

### 1.7.1 The basics

```
[37]: address = './Data/mtcars.csv'
cars = pd.read_csv(address)

cars.columns =
↳ ['car_names', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am', 'gear', 'carb']
cars.index = cars.car_names
cars.head(15)
```

```
[37]:
```

	car_names	mpg	cyl	disp	hp	drat	wt	\
car_names								
Mazda RX4	Mazda RX4	21.0	6	160.0	110	3.90	2.620	
Mazda RX4 Wag	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	
Datsun 710	Datsun 710	22.8	4	108.0	93	3.85	2.320	
Hornet 4 Drive	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	
Hornet Sportabout	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	
Valiant	Valiant	18.1	6	225.0	105	2.76	3.460	
Duster 360	Duster 360	14.3	8	360.0	245	3.21	3.570	
Merc 240D	Merc 240D	24.4	4	146.7	62	3.69	3.190	
Merc 230	Merc 230	22.8	4	140.8	95	3.92	3.150	
Merc 280	Merc 280	19.2	6	167.6	123	3.92	3.440	
Merc 280C	Merc 280C	17.8	6	167.6	123	3.92	3.440	
Merc 450SE	Merc 450SE	16.4	8	275.8	180	3.07	4.070	
Merc 450SL	Merc 450SL	17.3	8	275.8	180	3.07	3.730	
Merc 450SLC	Merc 450SLC	15.2	8	275.8	180	3.07	3.780	
Cadillac Fleetwood	Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	

	qsec	vs	am	gear	carb
car_names					
Mazda RX4	16.46	0	1	4	4
Mazda RX4 Wag	17.02	0	1	4	4



Datsun 710	18.61	1	1	4	1
Hornet 4 Drive	19.44	1	0	3	1
Hornet Sportabout	17.02	0	0	3	2
Valiant	20.22	1	0	3	1
Duster 360	15.84	0	0	3	4
Merc 240D	20.00	1	0	4	2
Merc 230	22.90	1	0	4	2
Merc 280	18.30	1	0	4	4
Merc 280C	18.90	1	0	4	4
Merc 450SE	17.40	0	0	3	3
Merc 450SL	17.60	0	0	3	3
Merc 450SLC	18.00	0	0	3	3
Cadillac Fleetwood	17.98	0	0	3	4

```
[38]: carb = cars.carb
      carb.value_counts()
```

```
[38]: 2    10
      4    10
      1     7
      3     3
      6     1
      8     1
      Name: carb, dtype: int64
```

```
[41]: cars_cat = cars[['cyl', 'vs', 'am', 'gear', 'carb']]
      cars_cat.head()
```

```
[41]:
```

	cyl	vs	am	gear	carb
car_names					
Mazda RX4	6	0	1	4	4
Mazda RX4 Wag	6	0	1	4	4
Datsun 710	4	1	1	4	1
Hornet 4 Drive	6	1	0	3	1
Hornet Sportabout	8	0	0	3	2

```
[114]: gears_group = cars_cat.groupby('gear')
      gears_group.describe()
```

```
[114]:
```

	vs								am								gear		carb \	
	count	mean	std	min	25%	50%	75%	max	count	mean	...	75%	max	count	mean					
cyl																				
4	11.0	0.9	0.3	0.0	1.0	1.0	1.0	1.0	11.0	0.7	...	4.0	5.0	11.0	1.5					
6	7.0	0.6	0.5	0.0	0.0	1.0	1.0	1.0	7.0	0.4	...	4.0	5.0	7.0	3.4					
8	14.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	14.0	0.1	...	3.0	5.0	14.0	3.5					

```

        std min 25% 50% 75% max
cyl
4    0.5 1.0 1.0 2.0 2.0 2.0
6    1.8 1.0 2.5 4.0 4.0 6.0
8    1.6 2.0 2.2 3.5 4.0 8.0

```

```
[3 rows x 32 columns]
```

### 1.7.2 Transforming variables to categorical data type

```
[43]: cars['group'] = pd.Series(cars.gear, dtype="category")
```

```
[44]: cars['group'].dtypes
```

```
[44]: CategoricalDtype(categories=[3, 4, 5], ordered=False)
```

```
[45]: cars['group'].value_counts()
```

```

[45]: 3    15
      4    12
      5     5
      Name: group, dtype: int64

```

### 1.7.3 Describing categorical data with crosstabs

```
[115]: pd.crosstab(cars['vs'], cars['cyl'])
```

```

[115]: cyl    4    6    8
      vs
      0     1    3   14
      1    10    4    0

```

## 1.8 Segment 5 - Starting with parametric methods in pandas and scipy

```

[47]: import matplotlib.pyplot as plt
      import seaborn as sb
      from pylab import rcParams

      import scipy
      from scipy.stats.stats import pearsonr

```

```

[48]: %matplotlib inline
      rcParams['figure.figsize'] = 8,4
      plt.style.use('seaborn-whitegrid')

```


### 1.8.1 The Pearson Correlation

```
[50]: address = './Data/mtcars.csv'

cars = pd.read_csv(address)
cars.columns = _
↳ ['car_names', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am', 'gear', 'carb']
```

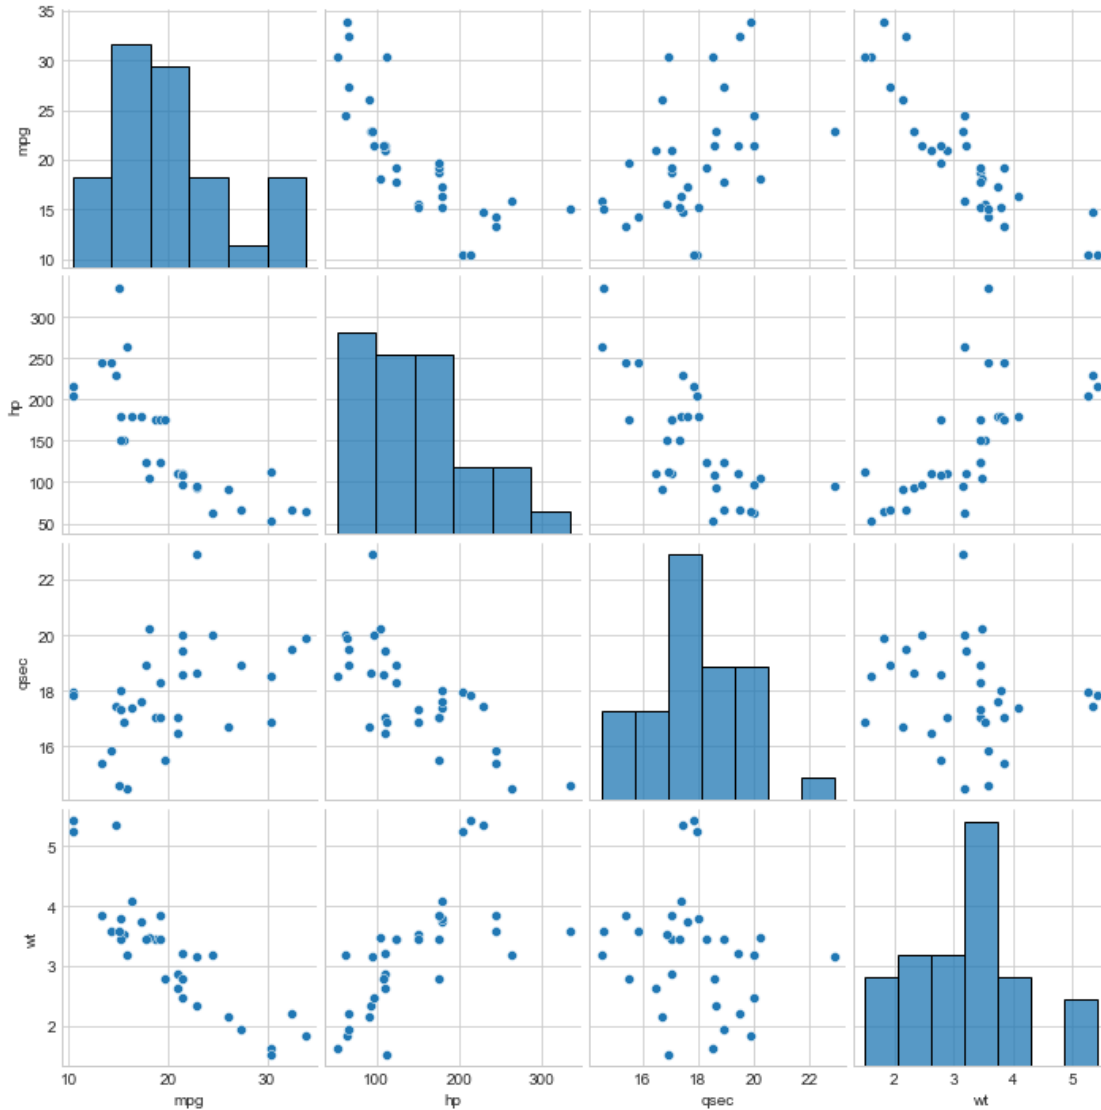
```
[51]: sb.pairplot(cars)
```

```
[51]: <seaborn.axisgrid.PairGrid at 0x1f9674c6610>
```



```
[52]: X = cars[['mpg', 'hp', 'qsec', 'wt']]
      sb.pairplot(X)
```

```
[52]: <seaborn.axisgrid.PairGrid at 0x1f96bebf80>
```



### 1.8.2 Using scipy to calculate the Pearson correlation coefficient

```
[53]: mpg = cars['mpg']
      hp = cars['hp']
      qsec = cars['qsec']
      wt = cars['wt']

      pearsonr_coefficient, p_value = pearsonr(mpg, hp)
```

```
print('PearsonR Correlation Coefficient %0.3f'% (pearsonr_coefficient))
```

PearsonR Correlation Coefficient -0.776

```
[54]: pearsonr_coefficient, p_value = pearsonr(mpg, qsec)
print('PearsonR Correlation Coefficient %0.3f'% (pearsonr_coefficient))
```

PearsonR Correlation Coefficient 0.419

```
[55]: pearsonr_coefficient, p_value = pearsonr(mpg, wt)
print('PearsonR Correlation Coefficient %0.3f'% (pearsonr_coefficient))
```

PearsonR Correlation Coefficient -0.868

### 1.8.3 Using pandas to calculate the Pearson correlation coefficient

```
[57]: corr = X.corr()
corr
```


```
[57]:
```

	mpg	hp	qsec	wt
mpg	1.000000	-0.776168	0.418684	-0.867659
hp	-0.776168	1.000000	-0.708223	0.658748
qsec	0.418684	-0.708223	1.000000	-0.174716
wt	-0.867659	0.658748	-0.174716	1.000000

### 1.8.4 Using Seaborn to visualize the Pearson correlation coefficient

```
[58]: sb.heatmap(corr, xticklabels=corr.columns.values, yticklabels= corr.columns.
→values)
```

```
[58]: <AxesSubplot:>
```



## 1.9 Segment 6 - Delving into non-parametric methods using pandas and scipy

```
[59]: import scipy
      from scipy.stats import spearmanr
```

```
[60]: %matplotlib inline
      rcParams['figure.figsize'] = 14, 7
      plt.style.use('seaborn-whitegrid')
```

### 1.9.1 The Spearman Rank Correlation

```
[62]: address = './Data/mtcars.csv'

      cars = pd.read_csv(address)
      cars.columns = ['car_names', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am', 'gear', 'carb']
```

```
[63]: cars.head()
```

```
[63]:
```


	car_names	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	\
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	

	carb
0	4
1	4
2	1
3	1
4	2


```
[64]: sb.pairplot(cars)
```

```
[64]: <seaborn.axisgrid.PairGrid at 0x1f96f2d5dc0>
```



```
[65]: X = cars[['cyl', 'vs', 'am', 'gear']]
      sb.pairplot(X)
```

```
[65]: <seaborn.axisgrid.PairGrid at 0x1f9747721c0>
```



```
[68]: cyl = cars['cyl']
      vs = cars['vs']
      am = cars['am']
      gear = cars['gear']
```

```
spearmanr_coefficient, p_value = spearmanr(cyl, vs)
print('Spearman Rank Correlation Coefficient %0.3f' % (spearmanr_coefficient))
```

Spearman Rank Correlation Coefficient -0.814

```
[69]: spearmanr_coefficient, p_value = spearmanr(cyl, am)
      print('Spearman Rank Correlation Coefficient %0.3f' % (spearmanr_coefficient))
```

Spearman Rank Correlation Coefficient -0.522



```
[70]: spearmanr_coefficient, p_value = spearmanr(cyl, gear)
print('Spearman Rank Correlation Coefficient %0.3f' % (spearmanr_coefficient))
```

Spearman Rank Correlation Coefficient -0.564

## 1.9.2 Chi-square test for independence

```
[72]: table = pd.crosstab(cyl, am)

from scipy.stats import chi2_contingency
chi2, p, dof, expected = chi2_contingency(table.values)
print('Chi-square statistic %0.3f p_value %0.3f' % (chi2, p))
```

Chi-square statistic 8.741 p\_value 0.013

```
[73]: table = pd.crosstab(cyl, vs)

from scipy.stats import chi2_contingency
chi2, p, dof, expected = chi2_contingency(table.values)
print('Chi-square statistic %0.3f p_value %0.3f' % (chi2, p))
```

Chi-square statistic 21.340 p\_value 0.000

```
[74]: table = pd.crosstab(cyl, gear)

from scipy.stats import chi2_contingency
chi2, p, dof, expected = chi2_contingency(table.values)
print('Chi-square statistic %0.3f p_value %0.3f' % (chi2, p))
```

Chi-square statistic 18.036 p\_value 0.001

## 1.10 Segment 7 - Transforming dataset distributions

```
[76]: import sklearn
from sklearn import preprocessing
from sklearn.preprocessing import scale
```

```
[77]: %matplotlib inline
rcParams['figure.figsize'] = 5, 4
sb.set_style('whitegrid')
```


### 1.10.1 Normalizing and transforming features with MinMaxScaler() and fit\_transform()

```
[78]: address = './Data/mtcars.csv'

cars = pd.read_csv(address)
cars.columns = ['car_names', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', '
↳ 'vs', 'am', 'gear', 'carb']
```

```
[79]: mpg = cars.mpg  
plt.plot(mpg)
```

```
[79]: [<matplotlib.lines.Line2D at 0x1f976aac730>]
```




```
[80]: cars[['mpg']].describe()
```

```
[80]:          mpg  
count  32.000000  
mean    20.090625  
std      6.026948  
min     10.400000  
25%     15.425000  
50%     19.200000  
75%     22.800000  
max     33.900000
```


```
[81]: mpg_matrix = mpg.values.reshape(-1,1)  
  
scaled = preprocessing.MinMaxScaler()  
  
scaled_mpg = scaled.fit_transform(mpg_matrix)  
plt.plot(scaled_mpg)
```

```
[81]: [<matplotlib.lines.Line2D at 0x1f976af6430>]
```



```
[82]: scaled = preprocessing.MinMaxScaler(feature_range=(0,10))  
  
scaled_mpg = scaled.fit_transform(mpg_matrix)  
plt.plot(scaled_mpg)
```


```
[82]: [<matplotlib.lines.Line2D at 0x1f976b2bcd0>]
```



### 1.10.2 Using scale() to scale your features


```
[83]: standardized_mpg = scale(mpg, axis=0, with_mean=False, with_std=False)
plt.plot(standardized_mpg)
```

```
[83]: [<matplotlib.lines.Line2D at 0x1f977b50820>]
```



```
[84]: standardized_mpg = scale(mpg)
plt.plot(standardized_mpg)
```

```
[84]: [<matplotlib.lines.Line2D at 0x1f977ba1ee0>]
```



```
[85]: address = './Data/iris.data.csv'
df = pd.read_csv(filepath_or_buffer=address, header=None, sep=',')

df.columns=['Sepal Length', 'Sepal Width', 'Petal Length', 'Petal Width', 'Species']
```

```
[86]: X = df.iloc[:,0:4].values
y = df.iloc[:,4].values
df[:5]
```


```
[86]:   Sepal Length  Sepal Width  Petal Length  Petal Width Species
0           5.1           3.5           1.4           0.2  setosa
1           4.9           3.0           1.4           0.2  setosa
2           4.7           3.2           1.3           0.2  setosa
3           4.6           3.1           1.5           0.2  setosa
4           5.0           3.6           1.4           0.2  setosa
```

## 1.11 Segment 8 - Extreme value analysis using univariate methods

### 1.11.1 Identifying outliers from Tukey boxplots

```
[88]: df.boxplot(return_type='dict')
plt.plot()
```

```
[88]: []
```



```
[89]: Sepal_Width = X[:,1]
iris_outliers = (Sepal_Width > 4)
df[iris_outliers]
```

```
[89]:   Sepal Length  Sepal Width  Petal Length  Petal Width Species
15          5.7          4.4          1.5          0.4  setosa
32          5.2          4.1          1.5          0.1  setosa
33          5.5          4.2          1.4          0.2  setosa
```

```
[90]: Sepal_Width = X[:,1]
iris_outliers = (Sepal_Width < 2.05)
df[iris_outliers]
```

```
[90]:   Sepal Length  Sepal Width  Petal Length  Petal Width  Species
60          5.0          2.0          3.5          1.0  versicolor
```

### 1.11.2 Applying Tukey outlier labeling

```
[91]: pd.options.display.float_format = '{:.1f}'.format
X_df = pd.DataFrame(X)
print(X_df.describe())
```

```
      0      1      2      3
count 150.0 150.0 150.0 150.0
mean   5.8   3.1   3.8   1.2
std    0.8   0.4   1.8   0.8
min    4.3   2.0   1.0   0.1
```

25%	5.1	2.8	1.6	0.3
50%	5.8	3.0	4.3	1.3
75%	6.4	3.3	5.1	1.8
max	7.9	4.4	6.9	2.5


## 1.12 Segment 9 - Multivariate analysis for outlier detection

```
[94]: df = pd.read_csv(filepath_or_buffer='./Data/iris.data.csv', header=None,
    ↪sep=',')
```

```
df.columns=['Sepal Length', 'Sepal Width', 'Petal Length', 'Petal Width',
    ↪'Species']
```

```
[95]: data = df.iloc[:,0:4].values
target = df.iloc[:,4].values
df[:5]
sb.boxplot(x='Species', y='Sepal Length', data=df, palette='hls')
```


```
[95]: <AxesSubplot:xlabel='Species', ylabel='Sepal Length'>
```



### 1.12.1 Looking at the scatterplot matrix

```
[97]: sb.pairplot(df, hue='Species', palette='hls')
```

```
[97]: <seaborn.axisgrid.PairGrid at 0x1f977d736d0>
```



```
[99]: xx=np.array([[7.,9.],[5.,12.]])
      yy=np.array([[2.,8.],[7.,4.]])
      np.dot(xx,yy)
```

```
[99]: array([[77., 92.],
            [94., 88.]])
```

```
[102]: xx=np.array([[1.,2.,3.],[4.,5.,6.]])
      yy=np.array([[10.,11.],[20.,21.],[30.,31.]])
      np.dot(xx,yy)
```

```
[102]: array([[140., 146.],
            [320., 335.]])
```

```
[103]: a=np.array([1,8,2,6,3,8,5,5,5,5])
      b=np.array([17,16,20,18,22,15,21,15,17,22])
```



```
(a+b)/10
```

```
[103]: array([1.8, 2.4, 2.2, 2.4, 2.5, 2.3, 2.6, 2. , 2.2, 2.7])
```

```
[106]: a=np.array([10, 15, 20])  
b=[5, 7, 9]  
(a-b)*7
```

```
[106]: array([35, 56, 77])
```

```
[109]: Q1= 1.714  
Q3=1.936  
iqr =Q3-Q1  
1.75*(iqr)
```

```
[109]: 0.38849999999999996
```

```
[110]: 1.714 -0.38
```

```
[110]: 1.334
```

```
[111]: 1.936+0.39
```

```
[111]: 2.326
```

```
[ ]:
```